

[illegible]

First off

- Computers can't store letters, numbers, pictures, etc.
- Computers store bits: 1 or 0
- Given that, then how do computers interpret letter and number characters?

Encoding

- An encoding scheme is a set of rules that translates bits to letter and numbers characters
- Encode: convert into a coded form
- In ASCII each character is one byte (8 bits)
- 01100010 01101001 01110100 01110011
- b i t s

ASCII

- All English letters a-zA-Z0-9, punctuation, some things like space, line feed, tab, backspace, etc.
- 128 characters in total = 7 bits
- Read left to write converting bits to characters
- `01001000 01100101 01101100 01101100 01101111 00100000`
`01010111 01101111 01110010 01101100 01100100`
- Hello World

What about other languages?

- 128 characters covers English, but what about Korean, Chinese, Hindi, Arabic, Russian, etc.
- To encompass all of these characters and languages more than a single byte (8 bits) is needed
- BIG-5 is a double-byte encoding that covers Traditional Chinese characters
- GB18030 does the same thing but covers simplified and traditional Chinese

GB18030


bits	character
10000001 01000000	𠂇
10000001 01000001	𠂈
10000001 01000010	𠂉
10000001 01000011	𠂊
10000001 01000100	𠂋

Unicode Overview

- An attempt to unify all encoding standards
- A code table of 1,114,112 code points
- Enough to encode all European, Middle-Eastern, Far-Eastern, Southern, Northern, Western, pre-historic and future characters
- Big enough for unofficial private-use sections
- There is an unofficial section for Klingon
- How many bits does Unicode use to encode all of this?

none

Unicode

- Unicode is not an encoding
- Unicode is a table of code points for characters
- Does not concern itself with how to represent those code points as bits
- "65 stands for A, 66 stands for B and 9,731 stands for 
- There are several ways to encode Unicode code points into bits

UTF-X

- UTF-32 uses 32bits to encode all Unicode code points. Simple, but wastes space.
- UTF-16 and UTF-8 are variable-length encodings
 - (UTF-8) If a character can be encoded using a single byte it will
 - If it requires two bytes then it will use that instead
 - UTF-16 uses two bytes by default, up to 4 bytes
 - The encoding uses the highest bits to signal how many bytes the character uses

character encoding bits

A UTF-8 01000001

A UTF-16 00000000 01000001

A UTF-32 00000000 00000000 00000000 01000001

あ UTF-8 11100011 10000001 10000010

あ UTF-16 00110000 01000010

あ UTF-32 00000000 00000000 00110000 01000010

Code Points

- Characters are often referred to by their Unicode code point
- Written in hex to keep numbers short
- Starts with a U+
- $\text{A}_\circ = \text{U+1E00} = 7680^{\text{th}}$ character

Encoding Issues

bits	encoding	characters
11000100	01000010	Windows Latin 1 ÄB
11000100	01000010	Mac Roman fB
11000100	01000010	GB18030 腩

characters	encoding	bits
Föö	Windows Latin 1	01000110 11111000 11110110
Föö	Mac Roman	01000110 10111111 10011010
Föö	UTF-8	01000110 11000011 10111000 11000011 10110110


What encoding is this?

```
10000011 01000111 10000011 10010011 10000011 01010010 10000001 01011011
10000011 01100110 10000011 01000010 10000011 10010011 10000011 01001111
10000010 11001101 10010011 11101111 10000010 10110101 10000010 10101101
10000010 11001000 10000010 10100010
```

Well...

- Most of the bytes start with 1 so not ASCII
- Most is not valid UTF-8
- Mac Roman works but you get
 - ÉGÉìÉRÅ[ÉfÉBÉìÉOÇÕìÔÇμÇ≠Ç»Ç¢
- Its Japanese Shift-JIS
 - エンコーディングは難しくない
- Some document viewers (and browsers) will start to read the bits and guess what the encoding is



- There's also the "Unicode replacement character"  (U+FFFD)
- A program may decide to insert for any character it couldn't decode correctly when trying to handle Unicode
- If a document is saved with some characters gone or replaced, then those characters are really gone for good with no way to reverse-engineer them

Security Considerations

- Imagine a scenario where:
 - An input validation filter rejects characters such as <, >, ', and " in a Web-application accepting UTF-8 encoded text.
 - An attacker sends in a U+FF1C FULLWIDTH LESS-THAN SIGN < in place of the ASCII <.
 - The attacker's input looks like: <script>
 - After passing through the XSS filter unchanged, the input moves deeper into the application.
 - Another API, perhaps at the data access layer, is configured to use a different character set such as windows-1252.
 - On receiving the input, a data access layer converts the multi-byte UTF-8 text to the single-byte windows-1252 code page, forcing a best-fit conversion to the dangerous characters the original XSS filter was trying to block. 7.The attacker's input successfully persists to the database.

Best-fit Mappings

Target char	Target code point	Test code point	Name
o	\u006F	\u2134	SCRIPT SMALL O
o	\u006F	\u014D	LATIN SMALL LETTER O WITH MACRON
s	\u0073	\u017F	LATIN SMALL LETTER LONG S
l	\u0049	\u0131	LATIN SMALL LETTER DOTLESS I
i	\u0069	\u0129	LATIN SMALL LETTER I WITH TILDE
K	\u004B	\u212A	KELVIN SIGN
k	\u006B	\u0137	LATIN SMALL LETTER K WITH CEDILLA
A	\u0041	\uFF21	FULLWIDTH LATIN CAPITAL LETTER A
a	\u0061	\u03B1	GREEK SMALL LETTER ALPHA
"	\u0022	\u02BA	MODIFIER LETTER DOUBLE PRIME
"	\u0022	\u030E	COMBINING DOUBLE VERTICAL LINE ABOVE
"	\u0027	\uFF02	FULLWIDTH QUOTATION MARK
'	\u0027	\u02B9	MODIFIER LETTER PRIME
'	\u0027	\u030D	COMBINING VERTICAL LINE ABOVE
'	\u0027	\uFF07	FULLWIDTH APOSTROPHE
<	\u003C	\uFF1C	FULLWIDTH LESS-THAN SIGN
<	\u003C	\uFE64	SMALL LESS-THAN SIGN
<	\u003C	\u2329	LEFT-POINTING ANGLE BRACKET
<	\u003C	\u3008	LEFT ANGLE BRACKET
<	\u003C	\u00AB	LEFT-POINTING DOUBLE ANGLE QUOTATION MARK
>	\u003E	\u00BB	RIGHT-POINTING DOUBLE ANGLE QUOTATION MARK
>	\u003E	\u3009	RIGHT ANGLE BRACKET
>	\u003E	\u232A	RIGHT-POINTING ANGLE BRACKET
>	\u003E	\uFE65	SMALL GREATER-THAN SIGN
>	\u003E	\uFF1E	FULLWIDTH GREATER-THAN SIGN
:	\u003A	\u2236	RATIO
:	\u003A	\u0589	ARMENIAN FULL STOP
:	\u003A	\uFE13	PRESENTATION FORM FOR VERTICAL COLON
:	\u003A	\uFE55	SMALL COLON
:	\u003A	\uFF1A	FULLWIDTH COLON

Filter Bypass

- Overconsumption bug
 - ` " onerror="alert(1)"</ br>`
 - ` " onerror="alert(1)"</ br>`
- Character deletion
 - Unicode BOM (Byte Order Mark) U+FEFF
 - `<scr[U+FEFF]ipt>`
- String transformation
 - `toLowerCase("İ") == "i"`
 - `toLowerCase("scrİpt") == "script"`
 - Never assume: `len(x) != len(toLower(x))`
- Whitespace (assigned the whitespace category and whitespace binary property)
 - Ogham space mark U+1680
 - Mongolian vowel separator U+180E
 - ``

Other Attacks

- Buffer Overflow
- Subtle Crypto Bugs
- Phishing
 - Who had the | domain?

Resources

- <http://kunststube.net/encoding/>
- <http://www.joelonsoftware.com/articles/Unicode.html>
- <https://websec.github.io/unicode-security-guide/character-transformations/>
- <http://www.unicode.org/reports/tr36/>
- <http://www.unicode.org/reports/tr39/>
- <https://www.blackhat.com/presentations/bh-usa-09/WEBER/BHUSA09-Weber-UnicodeSecurityPreview-PAPER.pdf>